Problem Statement

1. Import the dataset and do usual exploratory analysis steps like checking the structure & characteristics of the dataset:

A. Data type of all columns in the "customers" table.

QUERY:

```
SELECT
   column_name,
   data_type
FROM
   `target.INFORMATION_SCHEMA.COLUMNS`
WHERE
   table_name = 'customers';
```

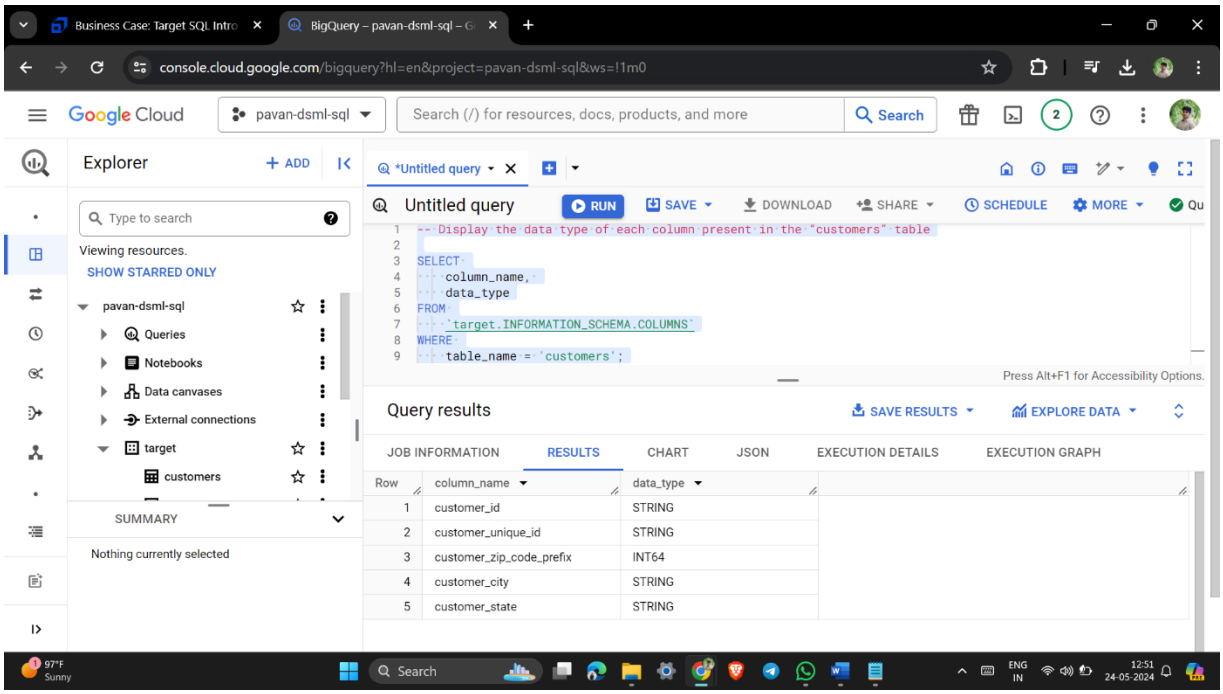

B. Get the time range between which the orders were placed.

QUERY:

```
SELECT
   MIN(order_purchase_timestamp) AS first_order_date,
   MAX(order_purchase_timestamp) AS last_order_date
```

```
FROM
    `target.orders`;
```

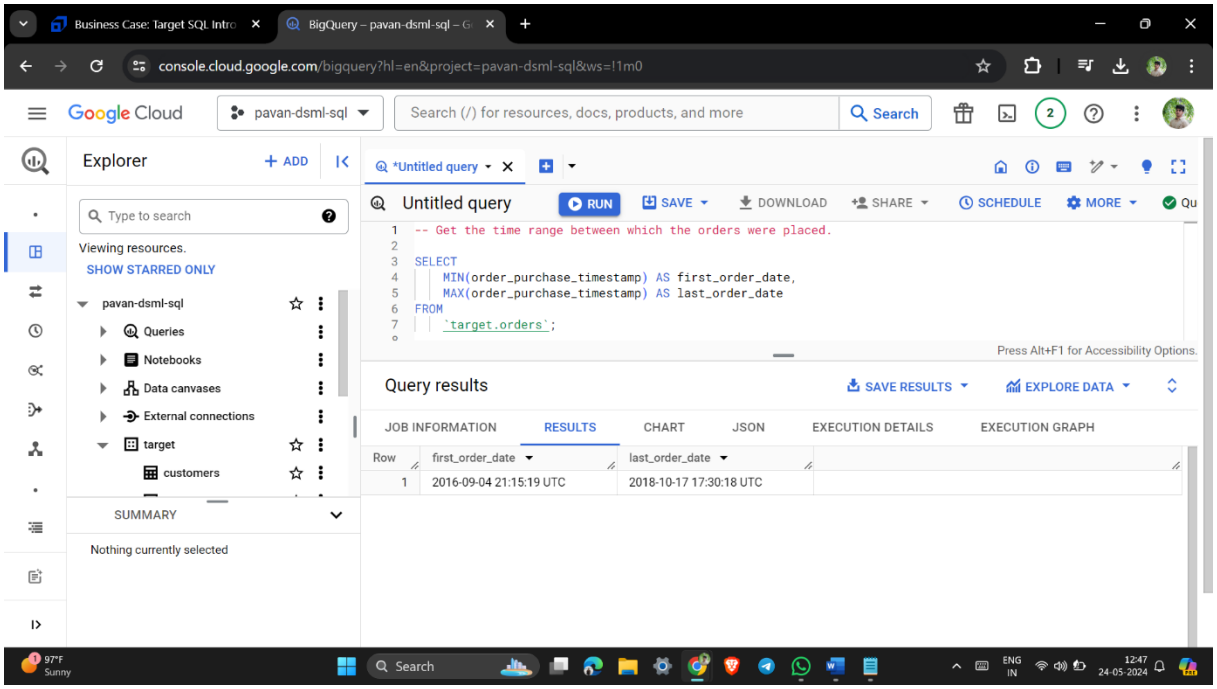

C. Count the Cities & States of customers who ordered during the given period.

QUERY:

```
SELECT
    COUNT(DISTINCT customer_city) AS unique_cities,
    COUNT(DISTINCT customer_state) AS unique_states
FROM
    `target.customers`;
```

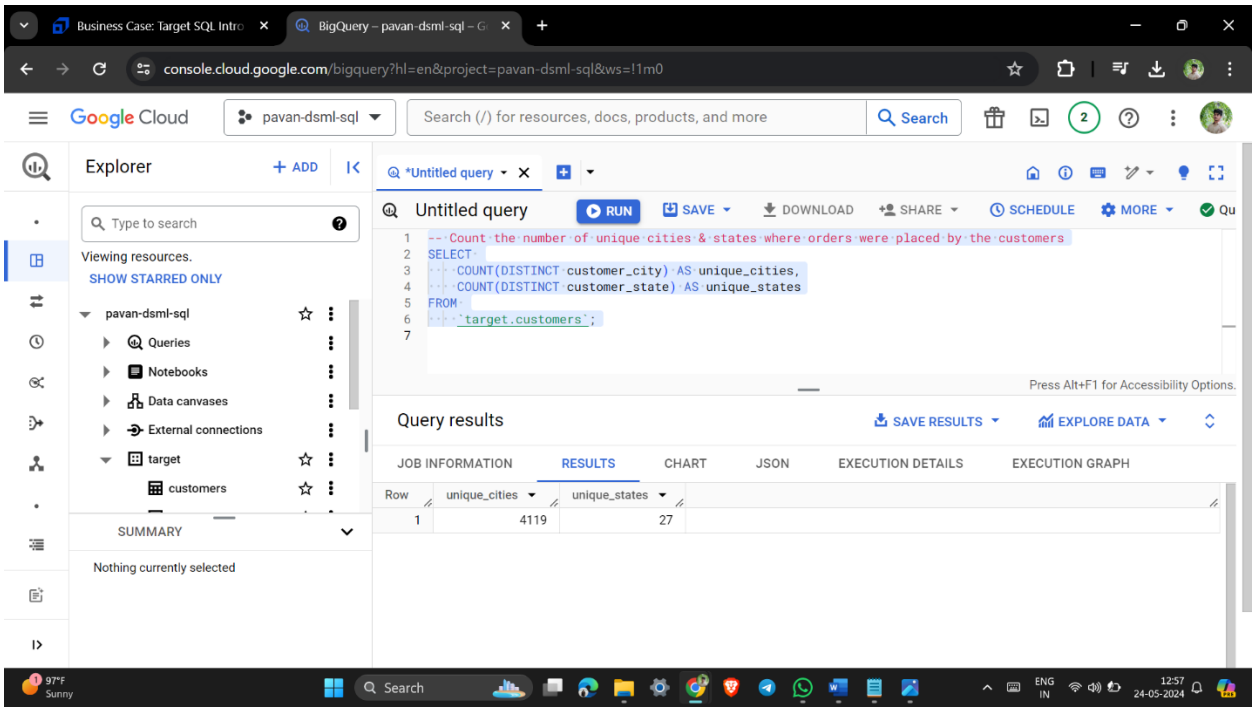

2. In-depthExploration:

A. Is there a growing trend in the no. of orders placed over the past years?

QUERY:

```
SELECT
  EXTRACT(YEAR FROM order_purchase_timestamp) AS year,
  EXTRACT(MONTH FROM order_purchase_timestamp) AS month,
  COUNT(order_id) AS number_of_orders
FROM
  `target.orders`
GROUP BY
  year, month
ORDER BY
  year, month;
```

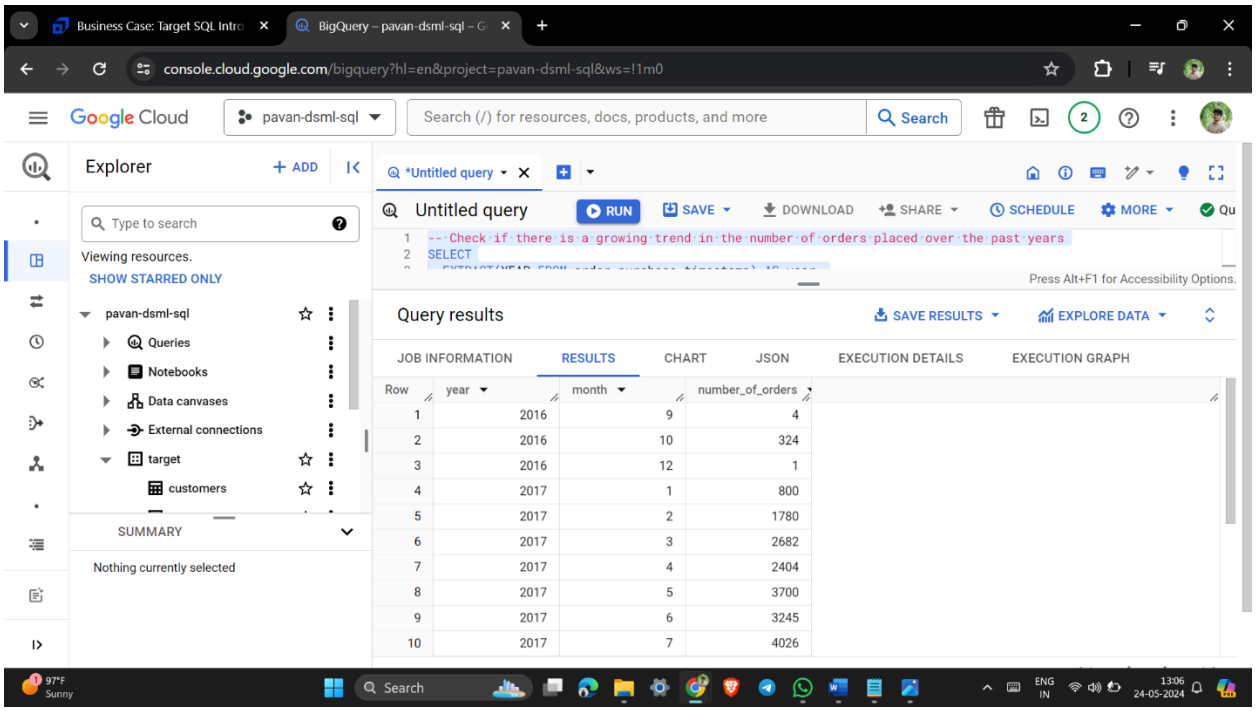

B. Can we see some kind of monthly seasonality in terms of the no. of orders being placed?

QUERY:

```
SELECT
  EXTRACT(MONTH FROM order_purchase_timestamp) AS month,
  COUNT(order_id) AS number_of_orders
FROM
  `target.orders`
GROUP BY
  month
ORDER BY
  month;
```

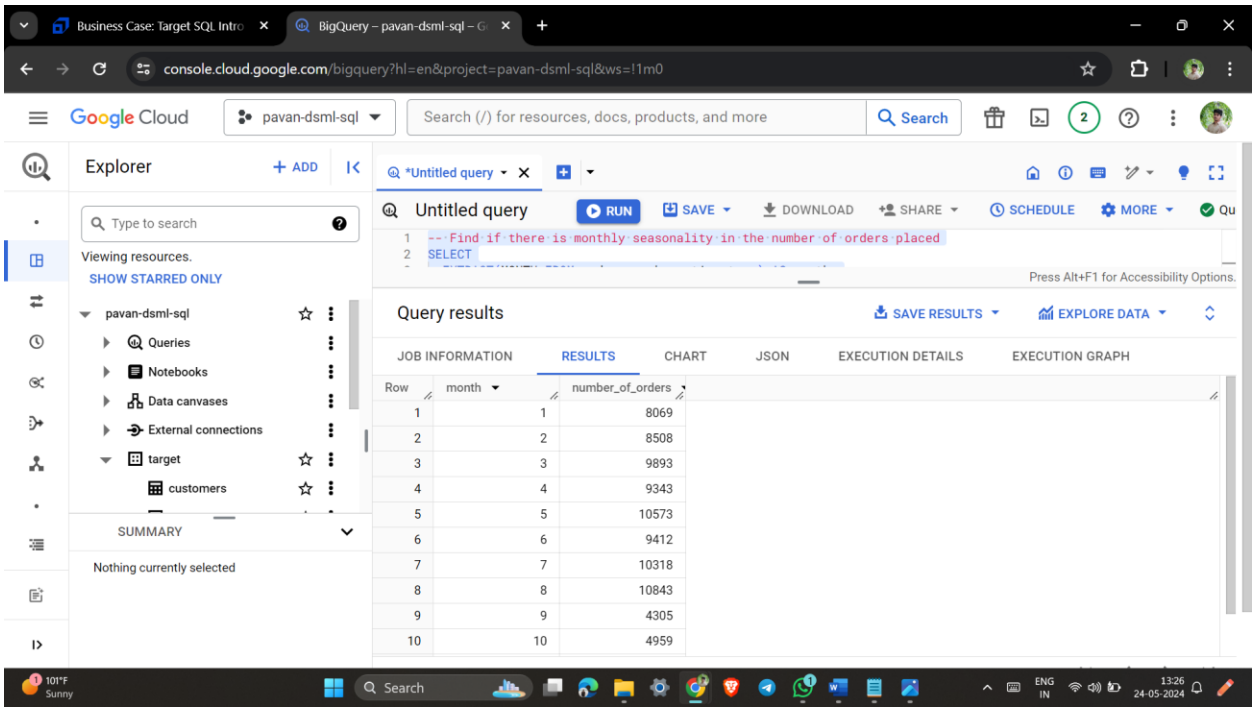

C. During what time of the day, do the Brazilian customers mostly place their orders? (Dawn, Morning, Afternoon or Night)

- **0-6 hrs : Dawn**
- **7-12 hrs : Mornings**
- **13-18 hrs : Afternoon**
- **19-23 hrs : Night**

QUERY:

```
SELECT
  CASE
    WHEN EXTRACT(HOUR FROM order_purchase_timestamp) BETWEEN 0
AND 6 THEN 'Dawn'
    WHEN EXTRACT(HOUR FROM order_purchase_timestamp) BETWEEN 7
AND 12 THEN 'Morning'
    WHEN EXTRACT(HOUR FROM order_purchase_timestamp) BETWEEN 13
AND 18 THEN 'Afternoon'
    WHEN EXTRACT(HOUR FROM order_purchase_timestamp) BETWEEN 19
AND 23 THEN 'Night'
  END AS time_of_day,
  COUNT(order_id) AS number_of_orders
FROM
  `target.orders`
GROUP BY
```

```
  time_of_day
ORDER BY
  number_of_orders DESC;
```

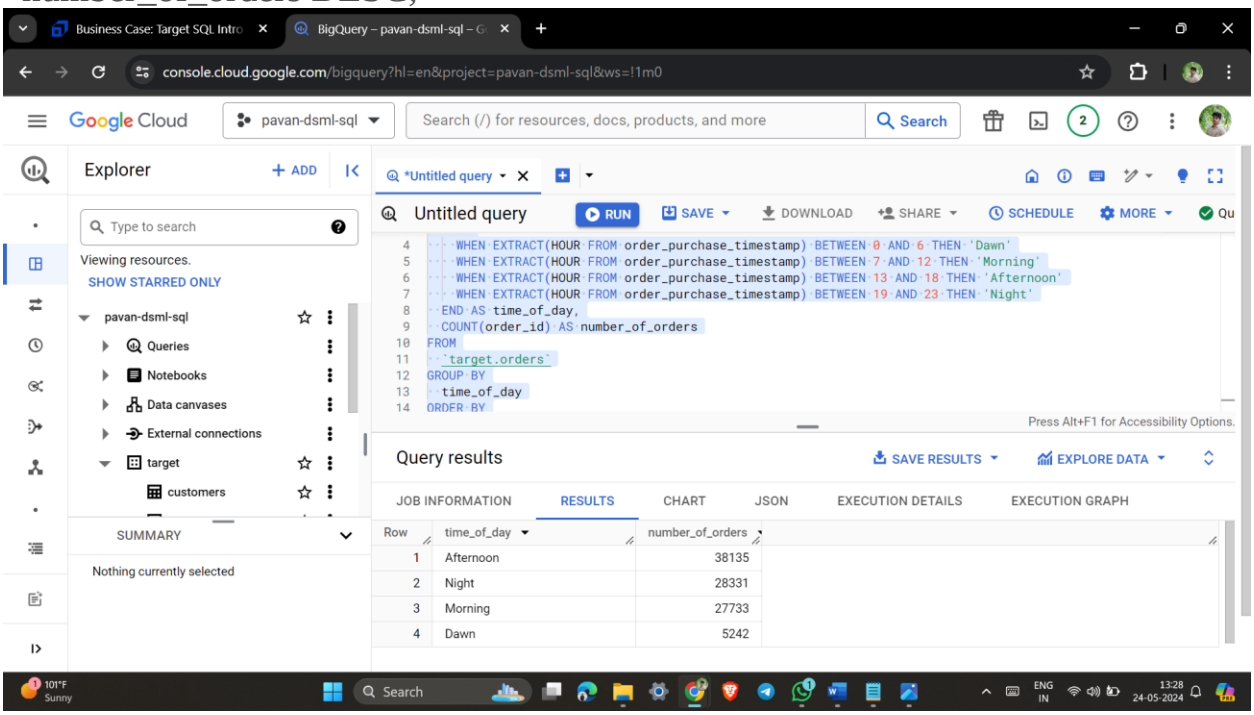

3. Evolution of E-commerce orders in the Brazil region:

A. Get the month-on-month no. of orders placed in each state.

QUERY:

```
  SELECT
    EXTRACT(YEAR FROM o.order_purchase_timestamp) AS year,
    EXTRACT(MONTH FROM o.order_purchase_timestamp) AS month,
    c.customer_state,
    COUNT(o.order_id) AS num_orders
  FROM
    `target.orders` AS o
  JOIN
    `target.customers` AS c
  ON
    o.customer_id = c.customer_id
  GROUP BY
    year, month, c.customer_state
  ORDER BY
```

year, month, c.customer_state;

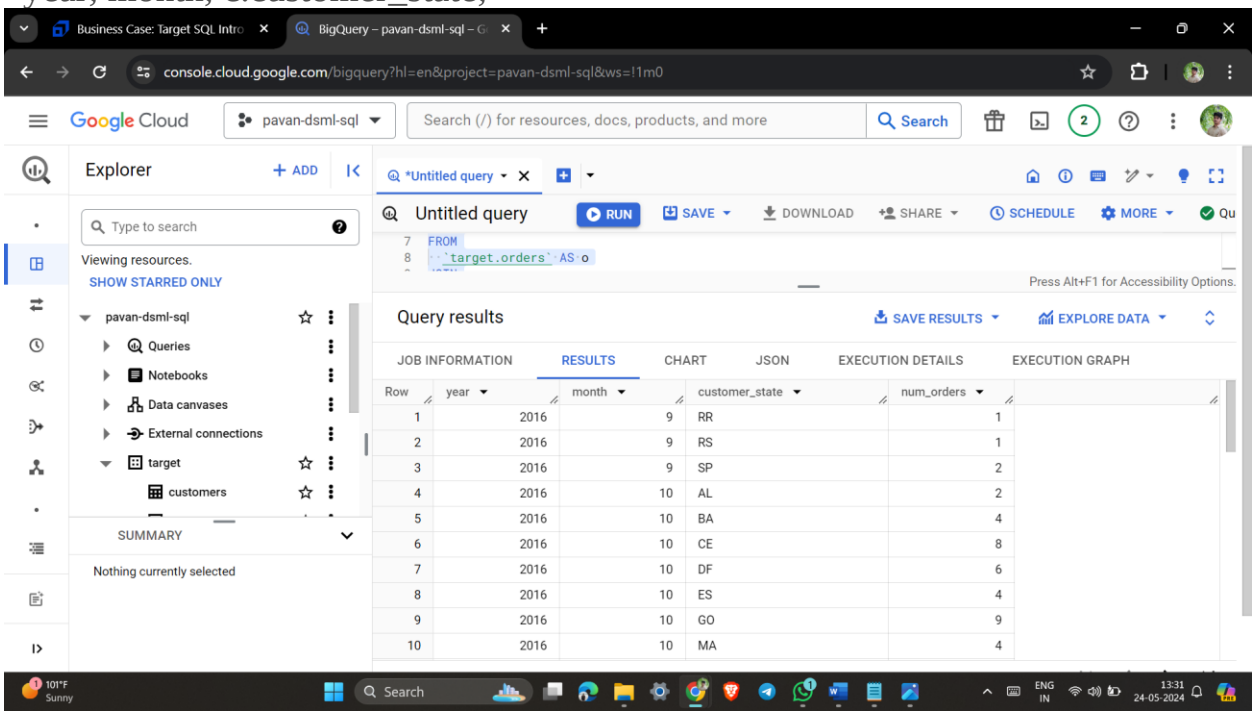

B. How are the customers distributed across all the states?

QUERY:

```
SELECT
  customer_state,
  COUNT(DISTINCT customer_id) AS num_customers
FROM
  `target.customers`
GROUP BY
  customer_state
ORDER BY
  customer_state;
```

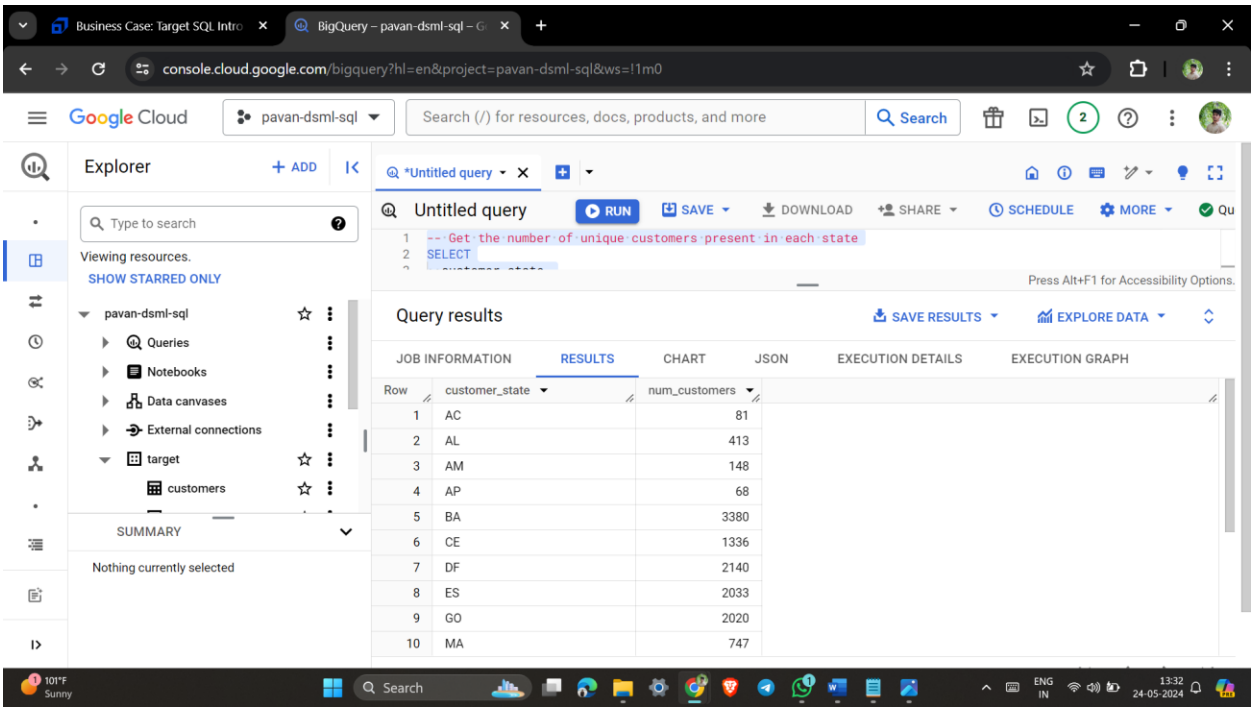

4. **Impact on Economy: Analyze the money movement by e-commerce by looking at order prices, freight and others.**

A. Get the % increase in the cost of orders from year 2017 to 2018 (include months between Jan to Aug only).
You can use the "payment_value" column in the payments table to get the cost of orders.

QUERY:

```
WITH order_costs AS (
  SELECT
    EXTRACT(YEAR FROM o.order_purchase_timestamp) AS year,
    EXTRACT(MONTH FROM o.order_purchase_timestamp) AS month,
    SUM(p.payment_value) AS total_payment_value
  FROM
    `target.orders` AS o
  JOIN
    `target.payments` AS p
  ON
    o.order_id = p.order_id
  WHERE
    EXTRACT(YEAR FROM o.order_purchase_timestamp) IN (2017, 2018)
    AND EXTRACT(MONTH FROM o.order_purchase_timestamp)
BETWEEN 1 AND 8
```

```
  GROUP BY
    year, month
)
SELECT
  (oc_2018.total_payment_value - oc_2017.total_payment_value) /
oc_2017.total_payment_value * 100 AS percent_increase
FROM
  order_costs oc_2017
JOIN
  order_costs oc_2018
ON
  oc_2017.month = oc_2018.month
WHERE
  oc_2017.year = 2017 AND oc_2018.year = 2018;
```

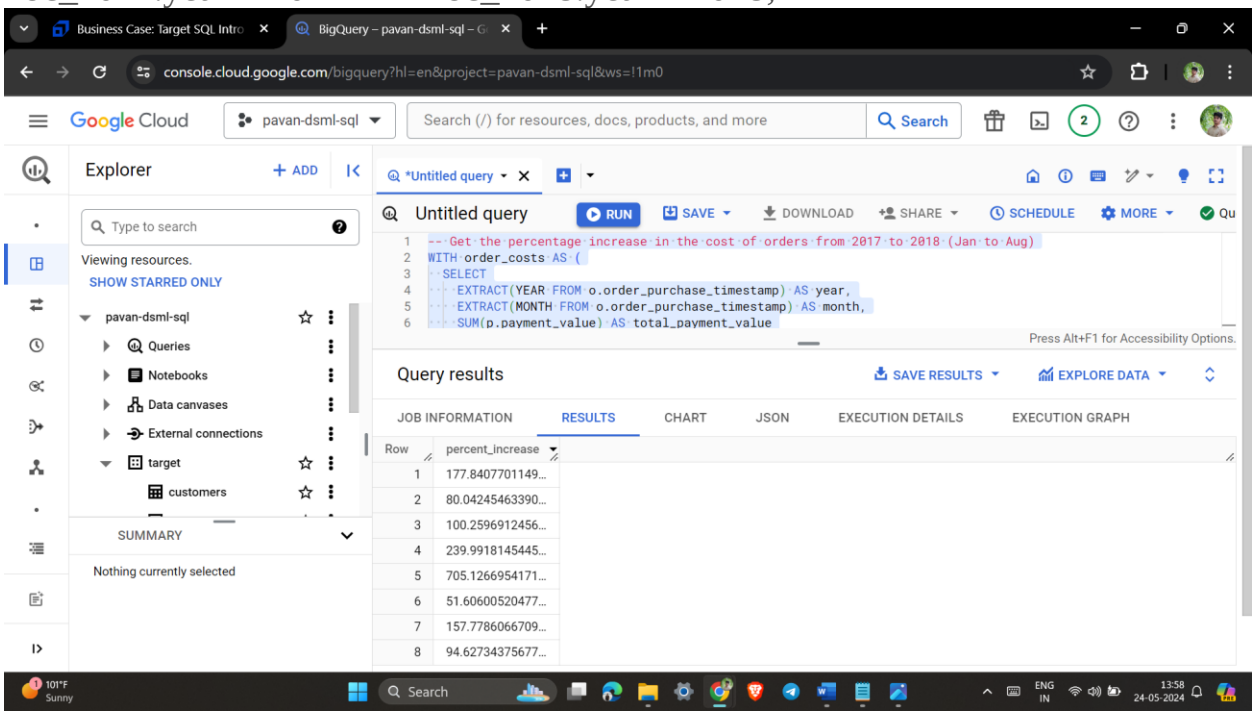

B. Calculate the Total & Average value of order price for each state.

QUERY:

```
SELECT
  c.customer_state,
  SUM(oi.price) AS total_order_price,
  AVG(oi.price) AS avg_order_price
FROM
  `target.order_items` AS oi
```

```
JOIN
  `target.orders` AS o
ON
  oi.order_id = o.order_id
JOIN
  `target.customers` AS c
ON
  o.customer_id = c.customer_id
GROUP BY
  c.customer_state
ORDER BY
  c.customer_state;
```

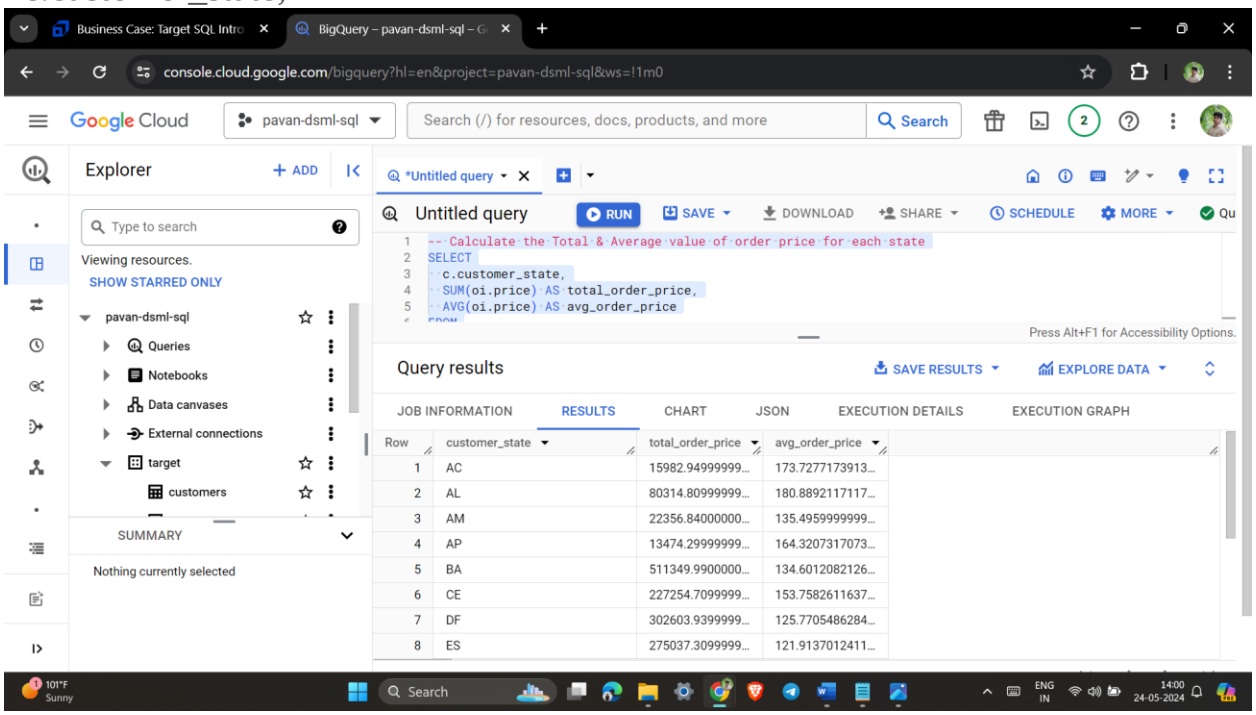

C. Calculate the Total & Average value of order freight for each state.

QUERY:

```
SELECT
  c.customer_state,
  SUM(oi.freight_value) AS total_freight_value,
  AVG(oi.freight_value) AS avg_freight_value
FROM
  `target.order_items` AS oi
JOIN
  `target.orders` AS o
```

```
ON
  oi.order_id = o.order_id
JOIN
  `target.customers` AS c
ON
  o.customer_id = c.customer_id
GROUP BY
  c.customer_state
ORDER BY
  c.customer_state;
```

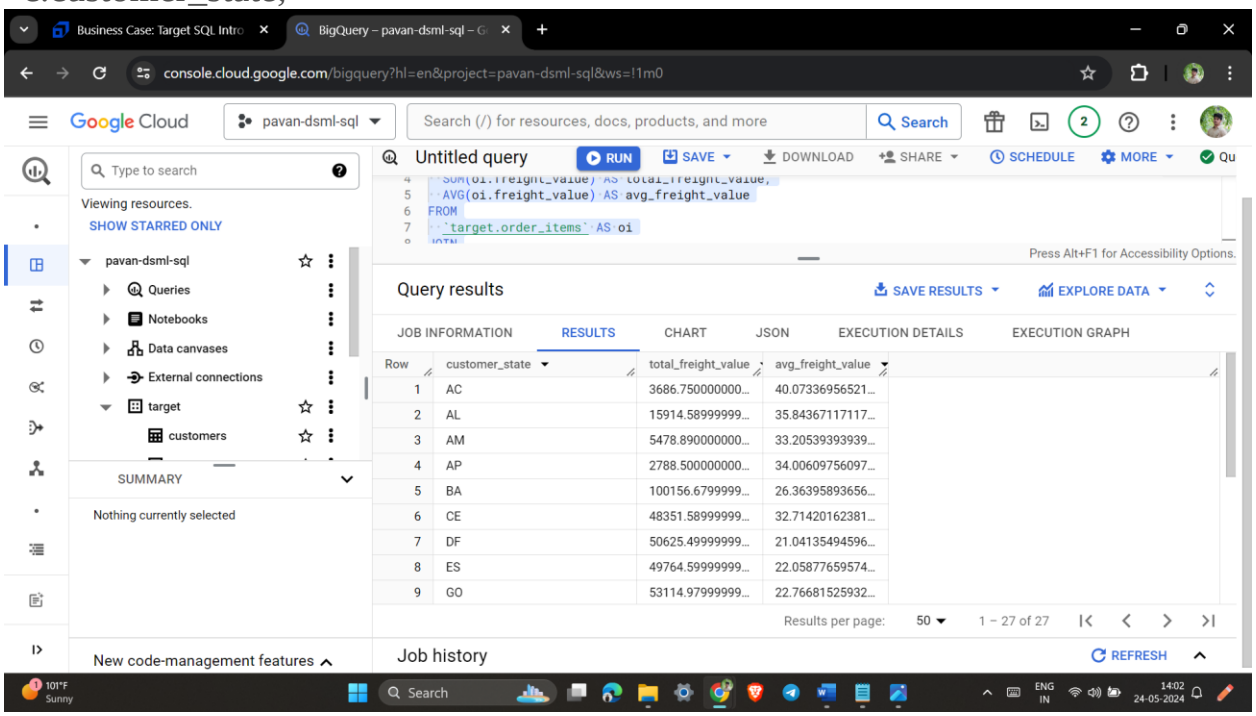

5. Analysis based on sales, freight and delivery time.

A. Find the no. of days taken to deliver each order from the order's purchase date as delivery time.
Also, calculate the difference (in days) between the estimated & actual delivery date of an order.
Do this in a single query.

You can calculate the delivery time and the difference between the estimated & actual delivery date using the given formula:

i. **time_to_deliver = order_delivered_customer_date - order_purchase_timestamp**

ii. **diff_estimated_delivery = order_delivered_customer_date - order_estimated_delivery_date**

QUERY:

```
SELECT
  order_id,
  DATE_DIFF(order_delivered_customer_date, order_purchase_timestamp,
DAY) AS delivery_time,
  DATE_DIFF(order_estimated_delivery_date, order_delivered_customer_date,
DAY) AS diff_estimated_delivery
FROM
  `target-project.target_brazil.orders`
WHERE
  order_status = 'delivered';
```

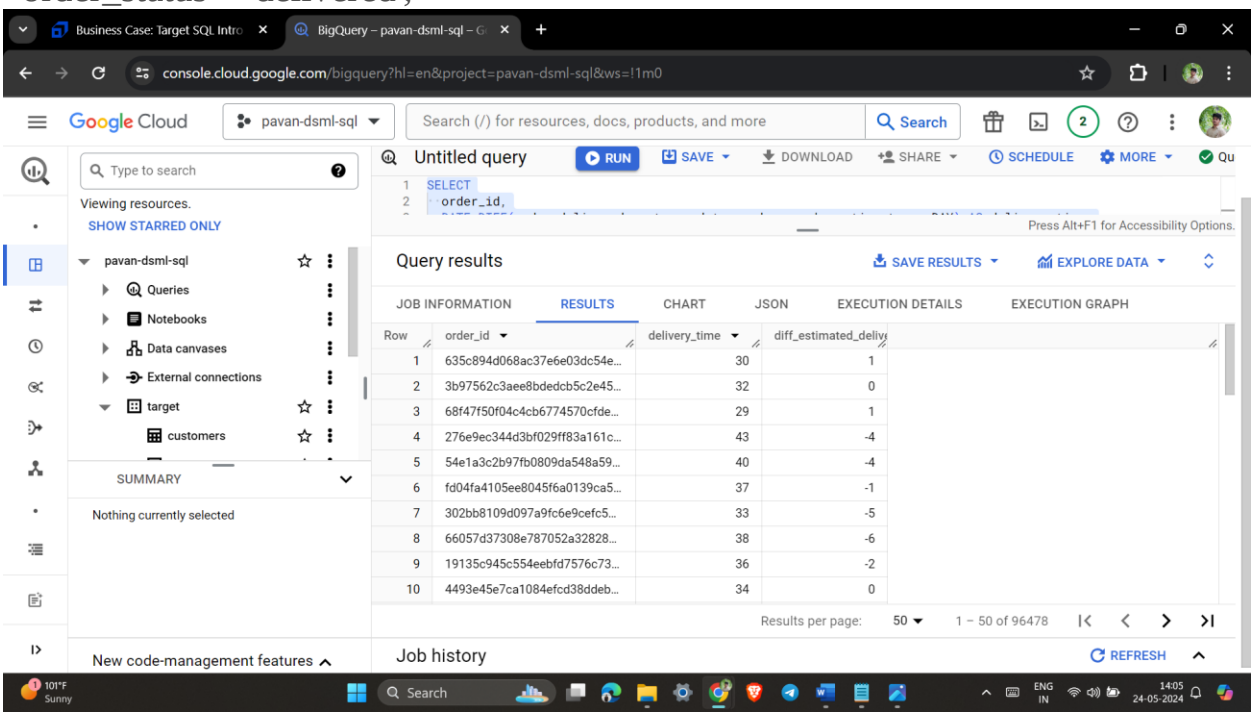

B. Find out the top 5 states with the highest & lowest average freight value.

QUERY:

```
with highest_freight as(
select customer_state,avg(freight_value) high_av,
row_number() over(order by avg(freight_value) desc) rhf
```

```
from `target.order_items` oi join `target.orders` o on oi.order_id=o.order_id
join `target.customers` c on c.customer_id=o.customer_id
group by customer_state
order by high_av desc
limit 5),
lowest_freight as(
select customer_state,avg(freight_value) low_av,
row_number() over(order by avg(freight_value) asc) rlf
from `target.order_items` oi join `target.orders` o on oi.order_id=o.order_id
join `target.customers` c on c.customer_id=o.customer_id
group by customer_state
order by low_av asc
limit 5)
select hf.customer_state,hf.high_av,lf.customer_state,lf.low_av from
highest_freight hf join lowest_freight lf
on hf.rhf=lf.rlf;
```

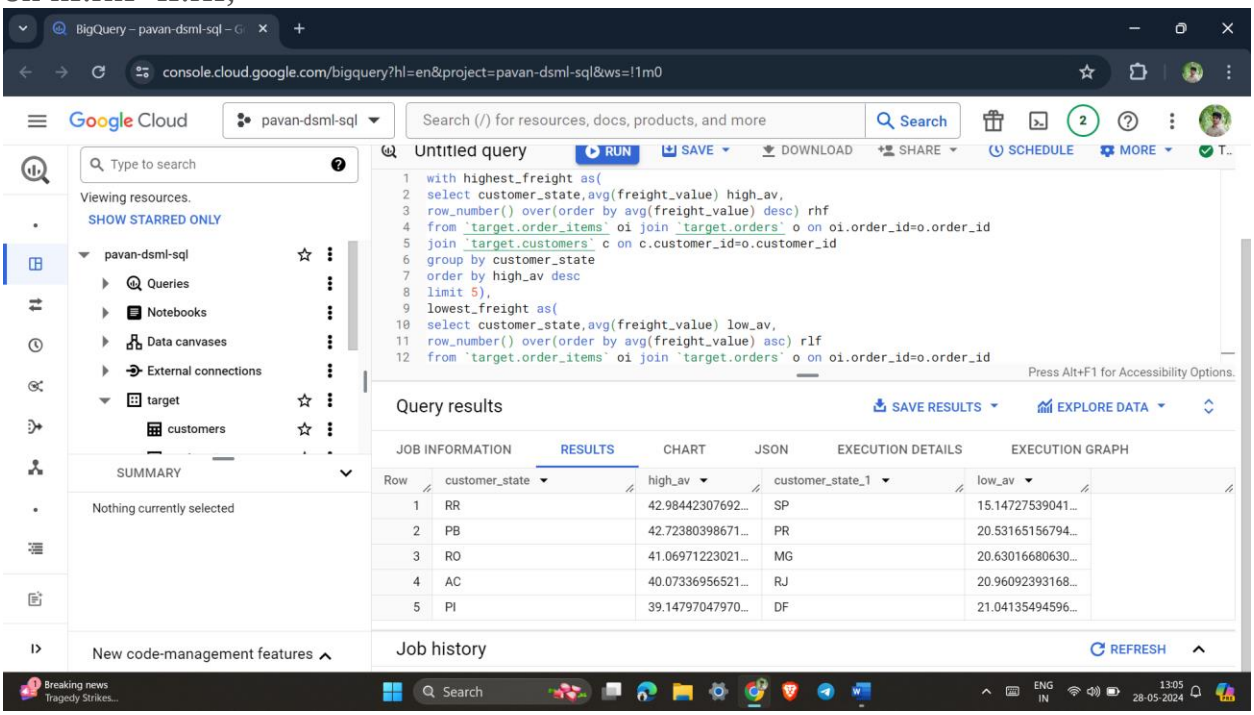

C. Find out the top 5 states with the highest & lowest average delivery time.

QUERY:

```
with cte as(select
customer_state,
```

```
avg(datetime_diff(order_delivered_customer_date,order_purchase_timesta
mp,day ) ) av_del_time,
from
`target.orders`o join  `target.customers`c on
o.customer_id=c.customer_id
group by customer_state),
cte2 as(
select *,
row_number() over(order by av_del_time desc) row_h,
from cte),
cte3 as(
  select *,
  row_number() over(order by av_del_time asc) row_l
  from cte
)
select
cte2.customer_state,cte2.av_del_time,cte3.customer_state,cte3.av_del_time
from
cte2 join cte3
on cte2.row_h=cte3.row_l
limit 5;
```

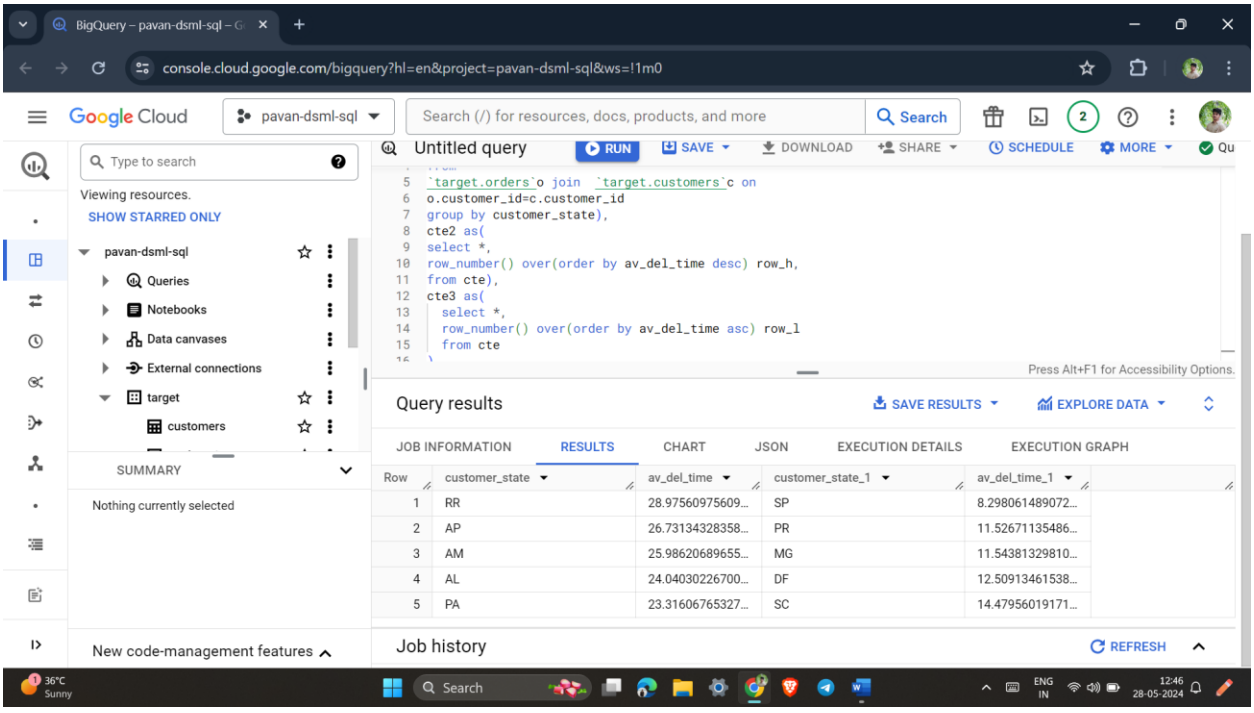

D. Find out the top 5 states where the order delivery is really fast as compared to the estimated date of delivery.
You can use the difference between the averages of actual & estimated delivery date to figure out how fast the delivery was for each state.

QUERY:

```
WITH fast_delivery_states AS (
  SELECT
    c.customer_state,
    AVG(DATE_DIFF(o.order_delivered_customer_date,
o.order_estimated_delivery_date, DAY)) AS avg_delivery_difference
  FROM
    `target.orders` AS o
  JOIN
    `target.customers` AS c
  ON
    o.customer_id = c.customer_id
  WHERE
    o.order_status = 'delivered'
  GROUP BY
    c.customer_state
)
SELECT
  customer_state,
  avg_delivery_difference
FROM
  fast_delivery_states
ORDER BY
  avg_delivery_difference ASC
LIMIT
  5;
```

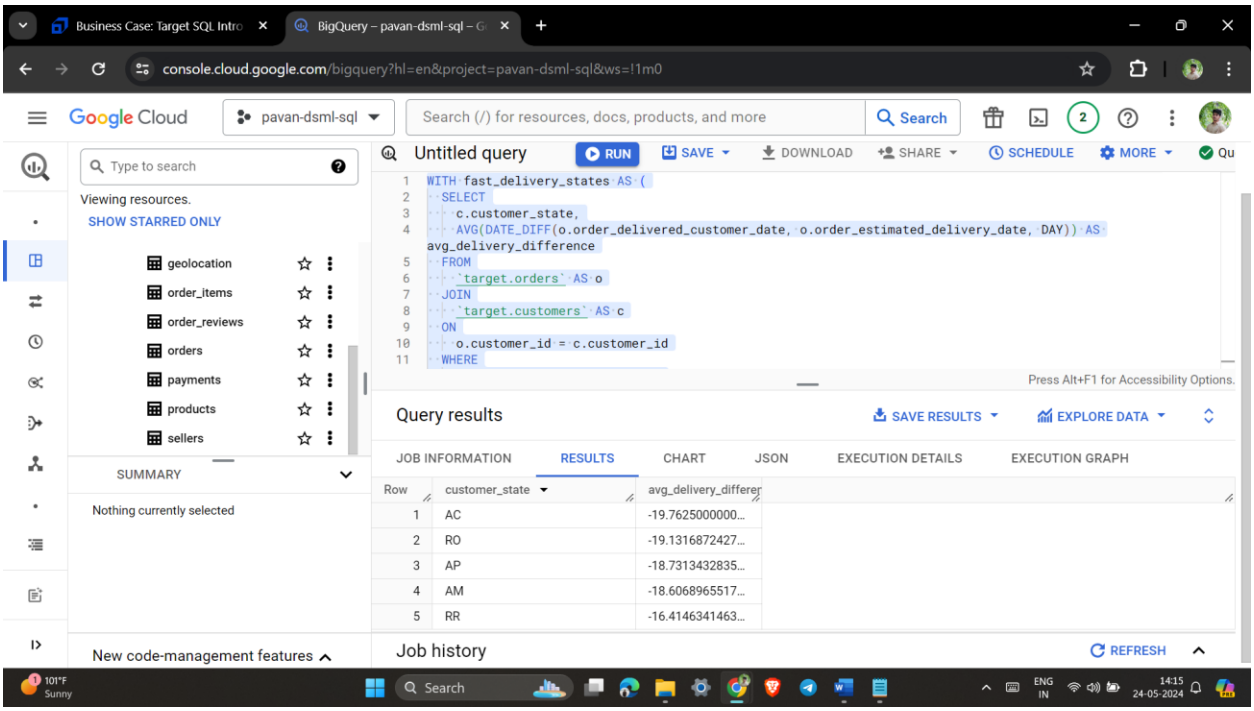

6. Analysis based on the payments:

A. Find the month-on-month no. of orders placed using different payment types.

QUERY:

```
SELECT
  EXTRACT(YEAR FROM o.order_purchase_timestamp) AS year,
  EXTRACT(MONTH FROM o.order_purchase_timestamp) AS month,
  p.payment_type,
  COUNT(DISTINCT o.order_id) AS num_orders
FROM
  `target.orders` AS o
INNER JOIN
  `target.payments` AS p
ON
  o.order_id = p.order_id
GROUP BY
  year, month, payment_type
ORDER BY
  year, month, payment_type;
```

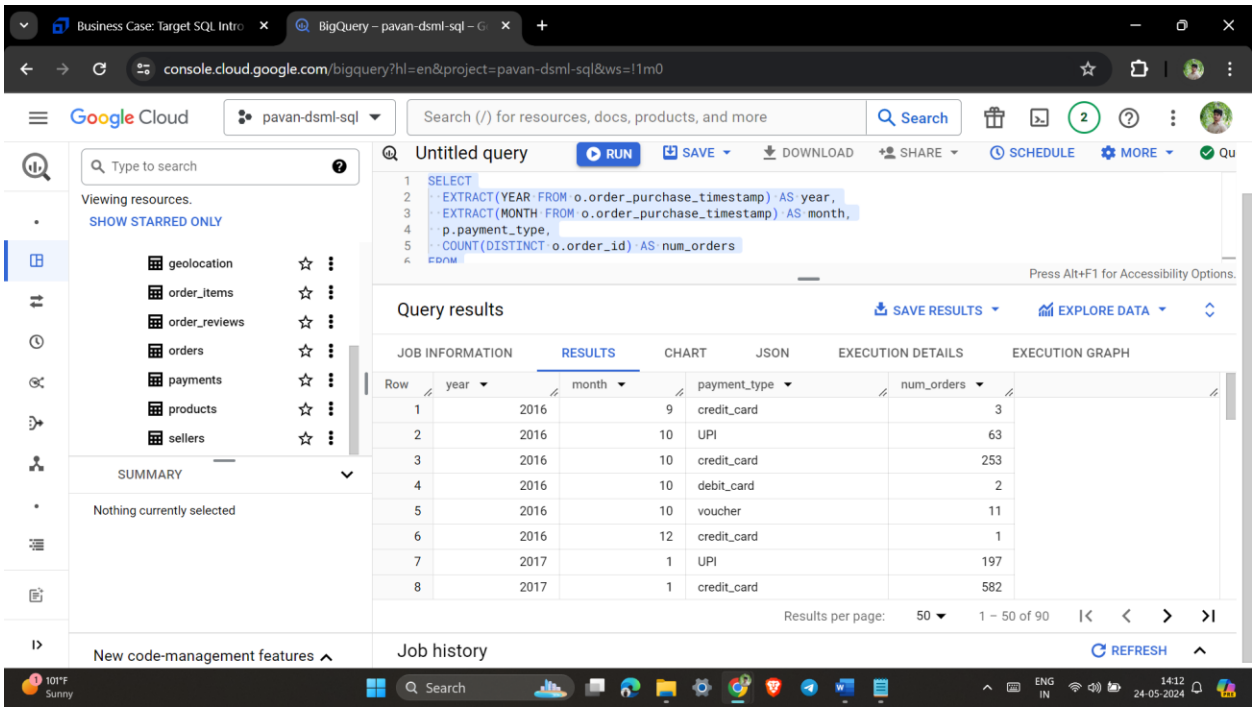

B. Find the no. of orders placed on the basis of the payment installments that have been paid.

QUERY:

```
SELECT
  payment_installments,
  COUNT(DISTINCT order_id) AS num_orders
FROM
  `target.payments`
WHERE
  payment_value > 0
GROUP BY
  payment_installments
ORDER BY
  payment_installments;
```

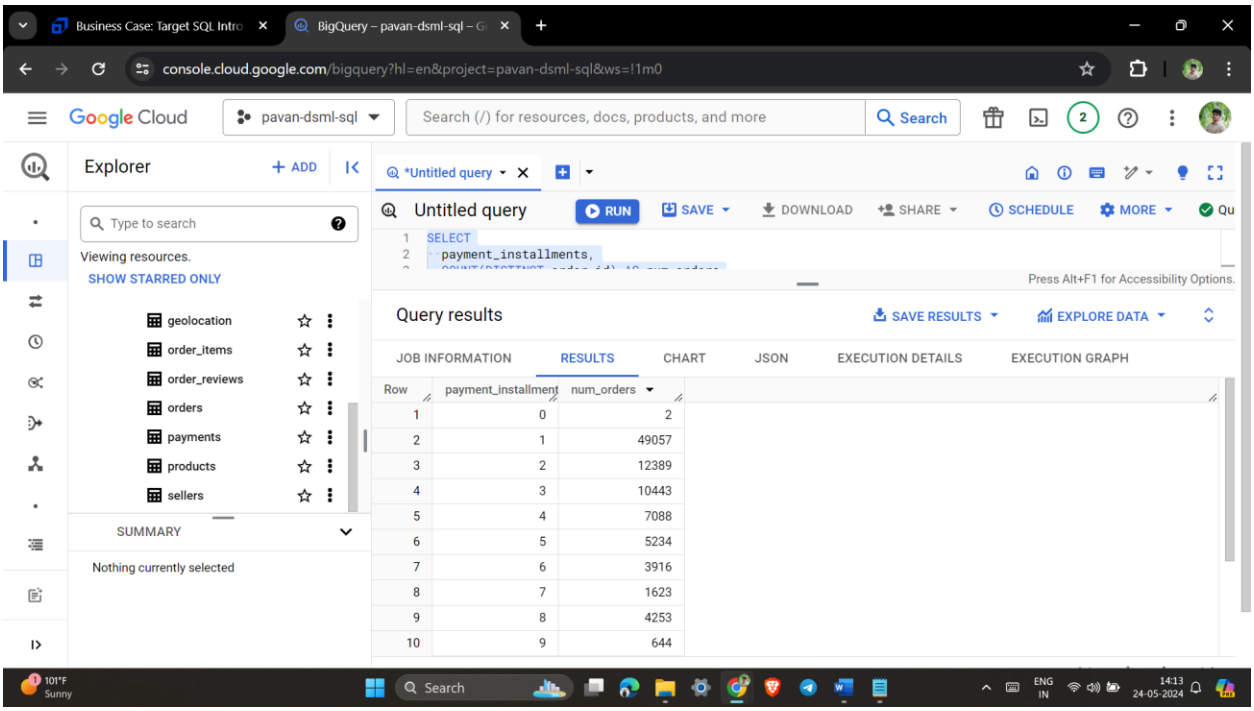
Business Case: Target SQL Intro
BigQuery – pavan-dsml-sql – G
console.cloud.google.com/bigquery?hl=en&project=pavan-dsml-sql&ws=!1m0
Google Cloud
pavan-dsml-sql
Search (/) for resources, docs, products, and more
Search
Explorer
+ ADD
Type to search
Viewing resources.
SHOW STARRED ONLY
geolocation
order_items
order_reviews
orders
payments
products
sellers
SUMMARY
Nothing currently selected
*Untitled query
Untitled query
RUN
SAVE
DOWNLOAD
SHARE
SCHEDULE
MORE
1 SELECT
2 payment_installments,
Press Alt+F1 for Accessibility Options.
Query results
SAVE RESULTS
EXPLORE DATA
JOB INFORMATION
RESULTS
CHART
JSON
EXECUTION DETAILS
EXECUTION GRAPH
Row
payment_installment
num_orders
1 0 2
2 1 49057
3 2 12389
4 3 10443
5 4 7088
6 5 5234
7 6 3916
8 7 1623
9 8 4253
10 9 644
101°F
Sunny
Search
ENG
IN
14:13
24-05-2024